

Using epoll in Linux

epoll is a Linux system call designed for efficiently monitoring multiple file descriptors for events, such as when data is available to read or a socket is ready to write. It is particularly useful in network programming, where a server might need to handle many connections simultaneously. Compared to older mechanisms like `select` or `poll`, epoll scales better with a large number of file descriptors because it avoids repeatedly scanning all descriptors for readiness.

Here's a step-by-step guide to using epoll in Linux, followed by a simple example of a server that accepts connections and echoes back data.

Steps to Use epoll

1. Create an Epoll Instance

To start using epoll, you need to create an epoll instance, which is represented by a file descriptor. Use the `epoll_create1` system call:

```
int epoll_fd = epoll_create1(0);
```

- **Parameter:** Pass 0 for basic usage (no special flags). Flags like `EPOLL_CLOEXEC` can be used for advanced scenarios.
- **Return Value:** Returns a file descriptor (`epoll_fd`) on success, or -1 on error (check `errno` for details).

The older `epoll_create` function is similar but takes a size hint (now ignored), so `epoll_create1` is preferred.

2. Add File Descriptors to Monitor

Use `epoll_ctl` to register file descriptors (e.g., sockets) with the epoll instance and specify the events you want to monitor:

```
struct epoll_event ev;  
ev.events = EPOLLIN; // Monitor for readability  
ev.data.fd = some_fd; // File descriptor to monitor  
epoll_ctl(epoll_fd, EPOLL_CTL_ADD, some_fd, &ev);
```

- **Parameters:**

- `epoll_fd`: The epoll instance file descriptor.
- `EPOLL_CTL_ADD`: Operation to add a file descriptor.

- `some_fd`: The file descriptor to monitor (e.g., a socket).
- `&ev`: Pointer to a `struct epoll_event` defining the events and optional user data.

- **Common Events:**

- `EPOLLIN`: Data available to read.
- `EPOLLOUT`: Ready to write.
- `EPOLLERR`: Error occurred.
- `EPOLLHUP`: Hang-up (e.g., connection closed).

- **User Data:** The `data` field in `struct epoll_event` can store a file descriptor (as shown) or other data (e.g., a pointer) to identify the source when events occur.

3. Wait for Events

Use `epoll_wait` to block and wait for events on the monitored file descriptors:

```
struct epoll_event events[MAX_EVENTS];
int nfds = epoll_wait(epoll_fd, events, MAX_EVENTS, -1);
```

- **Parameters:**

- `epoll_fd`: The epoll instance.
- `events`: Array to store triggered events.
- `MAX_EVENTS`: Maximum number of events to return (size of the array).
- `-1`: Timeout in milliseconds (-1 means wait indefinitely; 0 returns immediately).

- **Return Value:** Number of file descriptors with events (`nfds`), or -1 on error.

4. Handle Events

Loop through the events returned by `epoll_wait` and process them:

```
for (int i = 0; i < nfds; i++) {
    if (events[i].events & EPOLLIN) {
        // File descriptor events[i].data.fd is readable
    }
}
```

- Check the `events` field using bitwise operations (e.g., `events[i].events & EPOLLIN`) to determine the event type.
- Use `events[i].data.fd` to identify which file descriptor triggered the event.

5. Manage File Descriptors (Optional)

- **Remove:** Use `epoll_ctl` with `EPOLL_CTL_DEL` to stop monitoring a file descriptor:

```
epoll_ctl(epoll_fd, EPOLL_CTL_DEL, some_fd, NULL);
```

- **Modify:** Adjust events with `EPOLL_CTL_MOD`:

```
ev.events = EPOLLOUT; // Change to monitor writability  
epoll_ctl(epoll_fd, EPOLL_CTL_MOD, some_fd, &ev);
```

Key Concepts

Level-Triggered vs. Edge-Triggered

- **Level-Triggered (Default):** `epoll` notifies repeatedly as long as the condition persists (e.g., data remains unread). Simpler for most cases.
- **Edge-Triggered (EPOLLET):** Notifies only once when the state changes (e.g., new data arrives). Requires reading/writing all data until `EAGAIN` to avoid missing events; more efficient but trickier.
- Set `EPOLLET` in `ev.events` (e.g., `EPOLLIN | EPOLLET`) if using edge-triggered mode.

Non-Blocking I/O

`epoll` is often paired with non-blocking file descriptors to prevent blocking on I/O operations. Set a socket to non-blocking mode with:

```
fcntl(fd, F_SETFL, fcntl(fd, F_GETFL) | O_NONBLOCK);
```

Example: Simple Echo Server

Below is a basic example of a server that uses `epoll` to accept connections and echo data back to clients. It uses level-triggered mode for simplicity.

```
#include <sys/epoll.h>  
#include <sys/socket.h>  
#include <netinet/in.h>  
#include <fcntl.h>  
#include <unistd.h>
```

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>

#define MAX_EVENTS 10
#define PORT 8080

int main() {
    // Create listening socket
    int listen_fd = socket(AF_INET, SOCK_STREAM, 0);
    if (listen_fd == -1) { perror("socket"); exit(1); }

    struct sockaddr_in addr = { .sin_family = AF_INET, .sin_addr.s_addr = INADDR_ANY, .sin_port = htons(PORT) };
    if (bind(listen_fd, (struct sockaddr*)&addr, sizeof(addr)) == -1) { perror("bind"); exit(1); }
    if (listen(listen_fd, 5) == -1) { perror("listen"); exit(1); }

    // Set listening socket to non-blocking
    fcntl(listen_fd, F_SETFL, fcntl(listen_fd, F_GETFL) | O_NONBLOCK);

    // Create epoll instance
    int epoll_fd = epoll_create1(0);
    if (epoll_fd == -1) { perror("epoll_create1"); exit(1); }

    // Add listening socket to epoll
    struct epoll_event ev, events[MAX_EVENTS];
    ev.events = EPOLLIN; // Level-triggered
    ev.data.fd = listen_fd;
    if (epoll_ctl(epoll_fd, EPOLL_CTL_ADD, listen_fd, &ev) == -1) { perror("epoll_ctl"); exit(1); }

    // Event loop
    while (1) {
        int nfds = epoll_wait(epoll_fd, events, MAX_EVENTS, -1);
        if (nfds == -1) { perror("epoll_wait"); exit(1); }

        for (int i = 0; i < nfds; i++) {
            int fd = events[i].data.fd;

            if (fd == listen_fd) {
                // Accept new connection
                int client_fd = accept(listen_fd, NULL, NULL);

```

```

    if (client_fd == -1) { perror("accept"); continue; }

    // Set client socket to non-blocking
    fcntl(client_fd, F_SETFL, fcntl(client_fd, F_GETFL) | O_NONBLOCK);

    // Add client socket to epoll
    ev.events = EPOLLIN;
    ev.data.fd = client_fd;
    if (epoll_ctl(epoll_fd, EPOLL_CTL_ADD, client_fd, &ev) == -1) {
        perror("epoll_ctl: client_fd");
        close(client_fd);
    }
} else {
    // Handle client data
    char buf[1024];
    ssize_t n = read(fd, buf, sizeof(buf));
    if (n <= 0) {
        // Error or EOF: close and remove from epoll
        if (n == -1 && (errno == EAGAIN || errno == EWOULDBLOCK)) continue;
        close(fd);
        epoll_ctl(epoll_fd, EPOLL_CTL_DEL, fd, NULL);
    } else {
        // Echo data back
        write(fd, buf, n); // Note: In real apps, handle partial writes
    }
}
}

close(epoll_fd);
close(listen_fd);
return 0;
}

```

Notes on the Example

- **Listening Socket:** Monitors EPOLLIN to detect new connections.
- **Client Sockets:** Also monitor EPOLLIN to detect incoming data.
- **Simplification:** Assumes `write` completes fully. In production, buffer data and use EPOLLOUT for partial writes.

- **Error Handling:** Closes sockets on errors or EOF and removes them from epoll.
-

Summary

epoll provides an efficient way to handle multiple file descriptors in Linux: 1. Create an instance with `epoll_create1`. 2. Register file descriptors and events with `epoll_ctl`. 3. Wait for events with `epoll_wait`. 4. Process events in a loop, adjusting monitored events or removing descriptors as needed.

For simple applications, level-triggered mode is recommended. For high-performance needs, consider edge-triggered mode with careful handling of all available data. Always pair epoll with non-blocking I/O for best results.