

雲計算和大數據入門

這節課包含以下話題：

- Spark
- Hadoop
- Kubernetes
- Docker
- Flink
- MongoDB

說起雲計算，似乎離不開很多的工具，Hadoop、Hive、Hbase、ZooKeeper、Docker、Kubernetes、Spark、Kafka、MongoDB、Flink、Druid、Presto、Kylin、Elastic Search。都有聽過嗎。這些工具有些我是從大數據工程師、分佈式後端工程師的職位描述上找到的。這些都是高薪職位。我們試著把他們都安裝上，試著把玩兩下。## 初探 Spark

官網說，Spark 用來處理大規模數據的分析引擎。spark 就是一套庫。它似乎不像 Redis 那樣分成服務端和客戶端。spark 就是只在客戶端使用的。從官網下載了最新的版本，spark-3.1.1-bin-hadoop3.2.tar^o

```
$ tree . -L 1
```

```
.  
  LICENSE  
  NOTICE  
  R  
  README.md  
  RELEASE  
  bin  
  conf  
  data  
  examples  
  jars  
  kubernetes  
  licenses  
  python  
  sbin  
  yarn
```

```
11 directories, 4 files
```

似乎就是各語言編寫的一些分析庫。

同時官網說可以在 Python 上直接裝依賴庫。pip install pyspark

```
$ pip install pyspark
Collecting pyspark
  Downloading pyspark-3.1.1.tar.gz (212.3 MB)
    | 212.3 MB 14 kB/s
Collecting py4j==0.10.9
  Downloading py4j-0.10.9-py2.py3-none-any.whl (198 kB)
    | 198 kB 145 kB/s
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.1.1-py2.py3-none-any.whl size=212767604 sha256=0b8079e8...
  Stored in directory: /Users/lzw/Library/Caches/pip/wheels/23/bf/e9/9f3500437422e2ab82246f25a51ee480a4...
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9 pyspark-3.1.1
```

裝上了。

這會看官網，有些例子

```
./bin/run-example SparkPi 10
```

哦，原來可以運行剛剛下載的安裝包裡的程序。但出錯了。

```
$ ./bin/run-example SparkPi 10
21/03/11 00:06:15 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using
21/03/11 00:06:16 INFO ResourceUtils: No custom resources configured for spark.driver.
21/03/11 00:06:16 WARN Utils: Service 'sparkDriver' could not bind on a random free port. You may check
```

Spark is a fast and general processing engine compatible with Hadoop data. It can run in Hadoop clusters through YARN or Spark's standalone mode, and it can process data in HDFS, HBase, Cassandra, Hive, and any Hadoop InputFormat. It is designed to perform both batch processing (similar to MapReduce) and new workloads like streaming, interactive queries, and machine learning.

出現了好幾次 hadoop。谷歌了 spark depends hadoop 之後，找到這樣一段話。看來這依賴於 Hadoop 格式的數據。讓我們先研究 Hadoop。

Hadoop

簡單看了官網後。來安裝一下。

```
brew install hadoop
```

安裝的過程中，來了解一下。

The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. It is designed to scale up from single servers to thousands of machines, each offering local computation and storage. Rather than rely on hardware to deliver high-availability, the library itself is designed to detect and handle failures at the application layer, so delivering a highly-available service on top of a cluster of computers, each of which may be prone to failures.

就是說 Hadoop 是一套框架，來處理分佈式的數據集。這些數據集可能分部在很多計算機上。用很簡單的編程模型來處理。它是設計來從單一服務器擴展到千台機器的。與其依賴於硬件的高可用，這個庫則設計來在應用層就能檢查和處理錯誤。因此能將高可用的服務部署到集群中，雖然集群中的每台電腦都可能導致失敗。

```
$ brew install hadoop
Error:
  homebrew-core is a shallow clone.
  homebrew-cask is a shallow clone.

To `brew update`, first run:
  git -C /usr/local/Homebrew/Library/Taps/homebrew/homebrew-core fetch --unshallow
  git -C /usr/local/Homebrew/Library/Taps/homebrew/homebrew-cask fetch --unshallow

These commands may take a few minutes to run due to the large size of the repositories.
This restriction has been made on GitHub's request because updating shallow
clones is an extremely expensive operation due to the tree layout and traffic of
Homebrew/homebrew-core and Homebrew/homebrew-cask. We don't do this for you
automatically to avoid repeatedly performing an expensive unshallow operation in
CI systems (which should instead be fixed to not use shallow clones). Sorry for
```

```

the inconvenience!

==> Downloading https://homebrew.bintray.com/bottles/openjdk-15.0.1.big_sur.bottle.tar.gz
Already downloaded: /Users/lzw/Library/Caches/Homebrew/downloads/d1e3ece4af1d225bc2607eaa4ce85a873d2c6d...
==> Downloading https://www.apache.org/dyn/closer.lua?path=hadoop/common/hadoop-3.3.0/hadoop-3.3.0.tar...
Already downloaded: /Users/lzw/Library/Caches/Homebrew/downloads/764c6a0ea7352bb8bb505989feeee1b36dc628c...
==> Installing dependencies for hadoop: openjdk
==> Installing hadoop dependency: openjdk
==> Pouring openjdk-15.0.1.big_sur.bottle.tar.gz
==> Caveats

For the system Java wrappers to find this JDK, symlink it with
  sudo ln -sfn /usr/local/opt/openjdk/libexec/openjdk.jdk /Library/Java/JavaVirtualMachines/openjdk.jdk

openjdk is keg-only, which means it was not symlinked into /usr/local,
because it shadows the macOS `java` wrapper.

If you need to have openjdk first in your PATH run:
  echo 'export PATH="/usr/local/opt/openjdk/bin:$PATH"' >> /Users/lzw/.bash_profile

For compilers to find openjdk you may need to set:
  export CPPFLAGS="-I/usr/local/opt/openjdk/include"

==> Summary
  /usr/local/Cellar/openjdk/15.0.1: 614 files, 324.9MB
==> Installing hadoop
  /usr/local/Cellar/hadoop/3.3.0: 21,819 files, 954.7MB, built in 2 minutes 15 seconds
==> Upgrading 1 dependent:
  maven 3.3.3 -> 3.6.3_1
==> Upgrading maven 3.3.3 -> 3.6.3_1
==> Downloading https://www.apache.org/dyn/closer.lua?path=maven/maven-3/3.6.3/binaries/apache-maven-3.0...
==> Downloading from https://mirror.olnevhost.net/pub/apache/maven/maven-3/3.6.3/binaries/apache-maven-3.0...
#####
Error: The `brew link` step did not complete successfully
The formula built, but is not symlinked into /usr/local
Could not symlink bin/mvn
Target /usr/local/bin/mvn
is a symlink belonging to maven. You can unlink it:

```

```
brew unlink maven
```

To force the link and overwrite all conflicting files:

```
brew link --overwrite maven
```

To list all files that would be deleted:

```
brew link --overwrite --dry-run maven
```

Possible conflicting files are:

```
/usr/local/bin/mvn -> /usr/local/Cellar/maven/3.3.3/bin/mvn
/usr/local/bin/mvnDebug -> /usr/local/Cellar/maven/3.3.3/bin/mvnDebug
/usr/local/bin/mvnyjp -> /usr/local/Cellar/maven/3.3.3/bin/mvnyjp
==> Summary
/usr/local/Cellar/maven/3.6.3_1: 87 files, 10.7MB, built in 7 seconds
Removing: /usr/local/Cellar/maven/3.3.3... (92 files, 9MB)
==> Checking for dependents of upgraded formulae...
==> No broken dependents found!
==> Caveats
==> openjdk
```

For the system Java wrappers to find this JDK, symlink it with

```
sudo ln -sfn /usr/local/opt/openjdk/libexec/openjdk.jdk /Library/Java/JavaVirtualMachines/openjdk.jdk
```

openjdk is keg-only, which means it was not symlinked into /usr/local, because it shadows the macOS `java` wrapper.

If you need to have openjdk first in your PATH run:

```
echo 'export PATH="/usr/local/opt/openjdk/bin:$PATH"' >> /Users/lzw/.bash_profile
```

For compilers to find openjdk you may need to set:

```
export CPPFLAGS="-I/usr/local/opt/openjdk/include"
```

注意到 brew 的輸出日誌中 maven 沒有很好地被鏈接。接下來，進行強制鏈接到 3.6.3_1 版本。

```
brew link --overwrite maven
```

Hadoop 就安裝成功了。

Modules

The project includes these modules:

- **Hadoop Common:** The common utilities that support the other Hadoop modules.
- **Hadoop Distributed File System (HDFS™):** A distributed file system that provides high-throughput access to application data.
- **Hadoop YARN:** A framework for job scheduling and cluster resource management.
- **Hadoop MapReduce:** A YARN-based system for parallel processing of large data sets.
- **Hadoop Ozone:** An object store for Hadoop.

說有這些模塊。這會敲入 hadoop 出現了：

```
$ hadoop
Usage: hadoop [OPTIONS] SUBCOMMAND [SUBCOMMAND OPTIONS]
or      hadoop [OPTIONS] CLASSNAME [CLASSNAME OPTIONS]
where CLASSNAME is a user-provided Java class
```

OPTIONS is none or any of:

--config dir	Hadoop config directory
--debug	turn on shell script debug mode
--help	usage information
buildpaths	attempt to add class files from build tree
hostnames list[,of,host,names]	hosts to use in slave mode
hosts filename	list of hosts to use in slave mode
loglevel level	set the log4j level for this command
workers	turn on worker mode

SUBCOMMAND is one of:

Admin Commands:

```
daemonlog      get/set the log level for each daemon
```

Client Commands:

```
archive        create a Hadoop archive
checknative    check native Hadoop and compression libraries availability
```

```
classpath      prints the class path needed to get the Hadoop jar and the required libraries
conftest       validate configuration XML files
credential     interact with credential providers
distch         distributed metadata changer
distcp         copy file or directories recursively
dtutil         operations related to delegation tokens
envvars        display computed Hadoop environment variables
fs             run a generic filesystem user client
gridmix        submit a mix of synthetic job, modeling a profiled from production load
jar <jar>       run a jar file. NOTE: please use "yarn jar" to launch YARN applications, not this command
jnopath        prints the java.library.path
kdiag          Diagnose Kerberos Problems
kerbname       show auth_to_local principal conversion
key            manage keys via the KeyProvider
rumenfolder    scale a rumen input trace
rumentrace     convert logs into a rumen trace
s3guard         manage metadata on S3
trace           view and modify Hadoop tracing settings
version         print the version
```

Daemon Commands:

```
kms            run KMS, the Key Management Server
registrydns   run the registry DNS server
```

SUBCOMMAND may print help when invoked w/o parameters or with -h.

官網給了些例子。

```
$ mkdir input
$ cp etc/hadoop/*.xml input
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.2.jar grep input output 'dfs[a-z]*'
$ cat output/*
```

注意到有 share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.2.jar。這意味著也許有些樣例文件我們沒有得到。猜測用 Homebrew 安裝會沒有這些文件。我們從官網下載了安裝文件包。

```
$ tree . -L 1
```

```
LICENSE-binary  
LICENSE.txt  
NOTICE-binary  
NOTICE.txt  
README.txt  
bin  
etc  
include  
lib  
libexec  
licenses-binary  
sbin  
share
```

出現了 share 目錄。然而 Homebrew 真的沒有附加的這些文件嗎。找到 Homebrew 安裝的目錄。

```
$ type hadoop  
hadoop is /usr/local/bin/hadoop  
$ ls -alrt /usr/local/bin/hadoop  
lrwxr-xr-x 1 lzw admin 33 Mar 11 00:48 /usr/local/bin/hadoop -> ../../Cellar/hadoop/3.3.0/bin/hadoop  
$ cd /usr/local/Cellar/hadoop/3.3.0
```

這是在 /usr/local/Cellar/hadoop/3.3.0/libexec/share/hadoop 下打印的目錄樹

```
“ ‘shell $ tree . -L 2 . |—— client | |—— hadoop-client-api-3.3.0.jar | |—— hadoop-client-minicluster-3.3.0.jar | |—— hadoop-client-runtime-3.3.0.jar |—— common | |—— hadoop-common-3.3.0-tests.jar | |—— hadoop-common-3.3.0.jar | |—— hadoop-kms-3.3.0.jar | |—— hadoop-nfs-3.3.0.jar | |—— hadoop-registry-3.3.0.jar | |—— jdiff | |—— lib | |—— sources | |—— webapps |—— hdfs | |—— hadoop-hdfs-3.3.0-tests.jar | |—— hadoop-hdfs-3.3.0.jar | |—— hadoop-hdfs-client-3.3.0-tests.jar | |—— hadoop-hdfs-client-3.3.0.jar | |—— hadoop-hdfs-httpfs-3.3.0.jar | |—— hadoop-hdfs-native-client-3.3.0-tests.jar | |—— hadoop-hdfs-native-client-3.3.0.jar | |—— hadoop-hdfs-nfs-3.3.0.jar | |—— hadoop-hdfs-rbf-3.3.0-tests.jar | |—— hadoop-hdfs-rbf-3.3.0.jar | |—— jdiff | |—— lib | |—— sources | |—— webapps |—— mapreduce | |—— hadoop-mapreduce-client-app-3.3.0.jar | |—— hadoop-mapreduce-client-common-3.3.0.jar | |—— hadoop-mapreduce-client-core-3.3.0.jar | |—— hadoop-mapreduce-client-hs-3.3.0.jar | |—— hadoop-mapreduce-client-hs-plugins-3.3.0.jar | |—— hadoop-mapreduce-client-jobclient-3.3.0-tests.jar | |——
```

```
hadoop-mapreduce-client-jobclient-3.3.0.jar |   ├── hadoop-mapreduce-client-nativetask-3.3.0.jar  
|   ├── hadoop-mapreduce-client-shuffle-3.3.0.jar |   ├── hadoop-mapreduce-client-uploader-  
3.3.0.jar |   ├── hadoop-mapreduce-examples-3.3.0.jar |   ├── jdiff |   ├── lib-examples  
|   └── sources └── tools |   ├── dynamometer |   ├── lib |   └── resource
```